



Interactive Ray Tracing using KD-Tree in DFS manner

Jiayin Cao



Agenda




- The highlight of my work.
- What's ray tracing.
- How do I map ray tracing algorithm on CUDA hardware.
- The performance of the algorithm.

The highlight of my work

- Mapping ray tracing algorithm on CUDA hardware well , support texture , shadow , reflection , refraction and multiple lights.
- A depth first search algorithm in GPU kernel without hardware stack.
- Achieving interactive performance for complex scenes.

The way ray tracing works



In computer graphics, ray tracing is a technique for generating an image by tracing the path of light through pixels in an image plane.

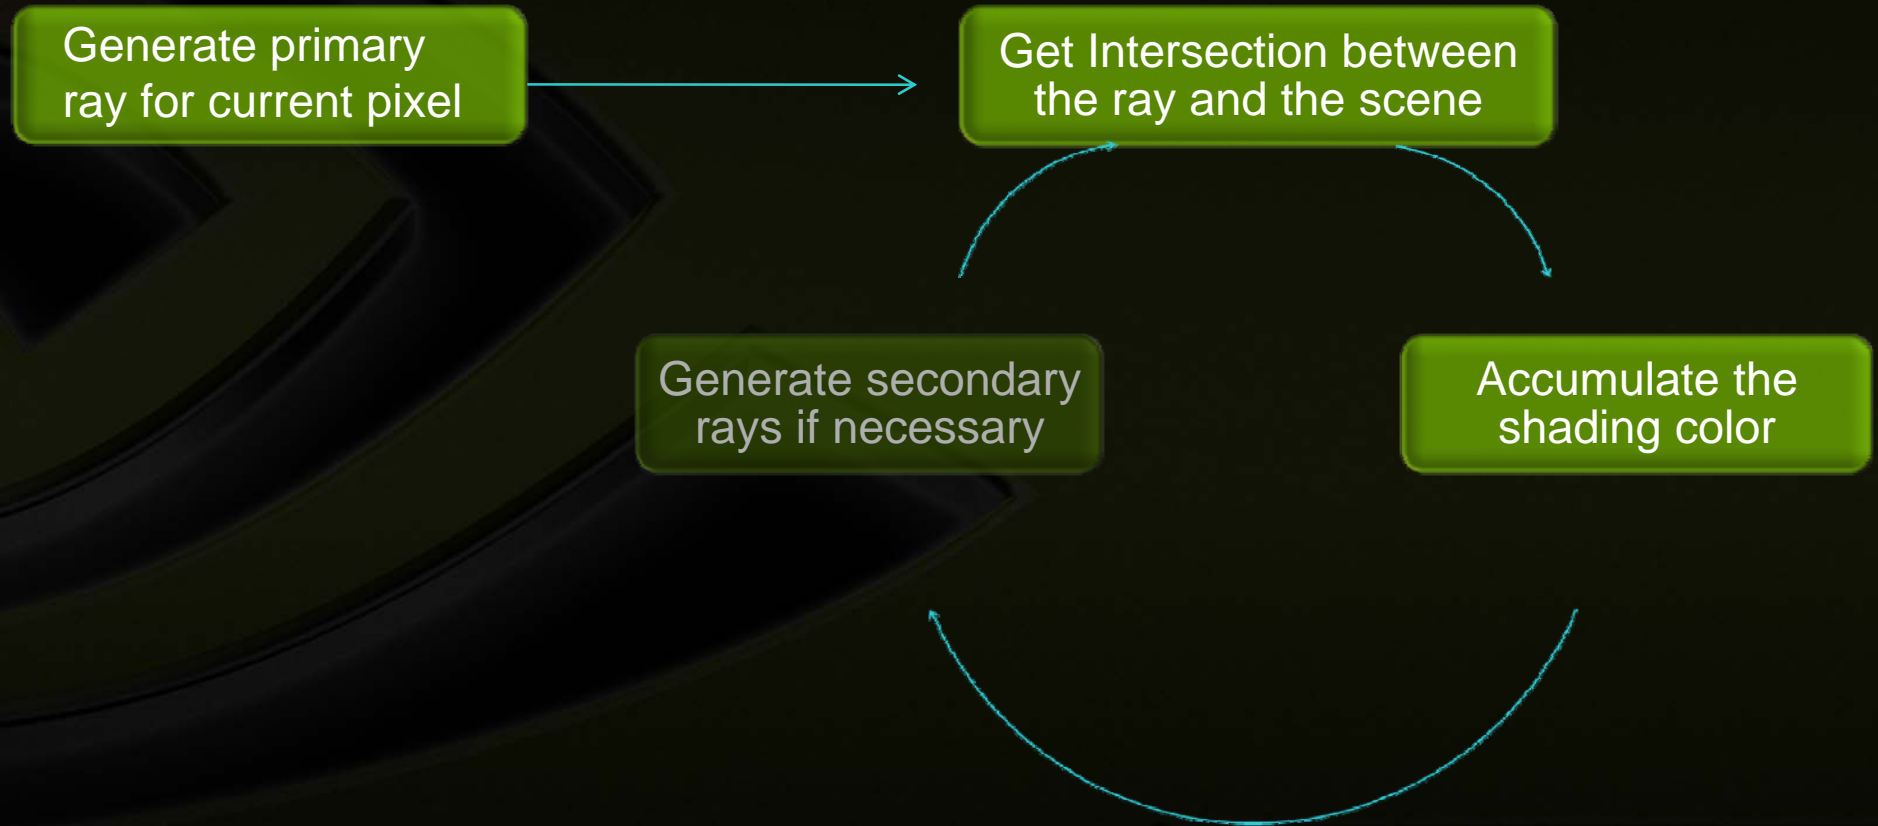
Why ray tracing

- Ray tracing can generate photo realistic images.(reflection, refraction)
- Ray tracing is software rendering , which means that every pass is programmable.
- Ray tracing could be mixed with other algorithm, such as photon mapping , radiosity.


A naïve ray tracing algorithm



- The following operations are performed for each pixel:



A naïve ray tracing algorithm



Accumulate
shading result
(get intersection ,
brdf)

Generate next
level ray


Generate
primary ray

The problems need to be solved



- How to map the algorithm on CUDA hardware.
- How to traverse KD-Tree to get the intersection result.

A simple way of doing GPU ray tracing



Each pixel is processed by one single thread.


Problem

Divergence in warp will cause bad performance.



thread



The way my algorithm works



Depth first search in KD-Tree

- **Problem:**
 - There is no hardware stack in GPU kernel.
- **The way it is solved:**
 - Using 32-bits unsigned integer as stack. (fast , low cost)



Depth first search in KD-Tree



KD-Tree Traverse

t → l
l → a
a → l
l → b
b → l
l → t (important)
t → r
r → d

Virtual stack using bits





Virtual stack

0
0
0
0
...

Traverse
state

Virtual stack using bits





Virtual stack

0
0
0
0
...

Traverse
state

Virtual stack using bits





Virtual stack

0
0
0
0
...

Traverse
state
→
Back track
state

Virtual stack using bits





Virtual stack

0
0
1
0
...

Traverse
state
→
Back track
state

Virtual stack using bits





Virtual stack

0
0
0
0
...

Back track
state

Virtual stack using bits





Virtual stack

0
0
0
0
...

Back track
state

Virtual stack using bits





Virtual stack

0
0
0
0
...

Traverse
state

Virtual stack using bits




Virtual stack

0
0
0
0
...

Traverse
state

The performance of the algorithm




Cornell Box (34 triangles)

GPU	FPS
GTX 285	30
GTX 460	50

70 times faster than a naïve CPU ray tracing algorithm.

The performance of the algorithm



Dragon on table (872870 triangles)

GPU	FPS
GTX 285	3
GTX 460	7

18 times faster than a naïve CPU ray tracing algorithm.

The gallery

